

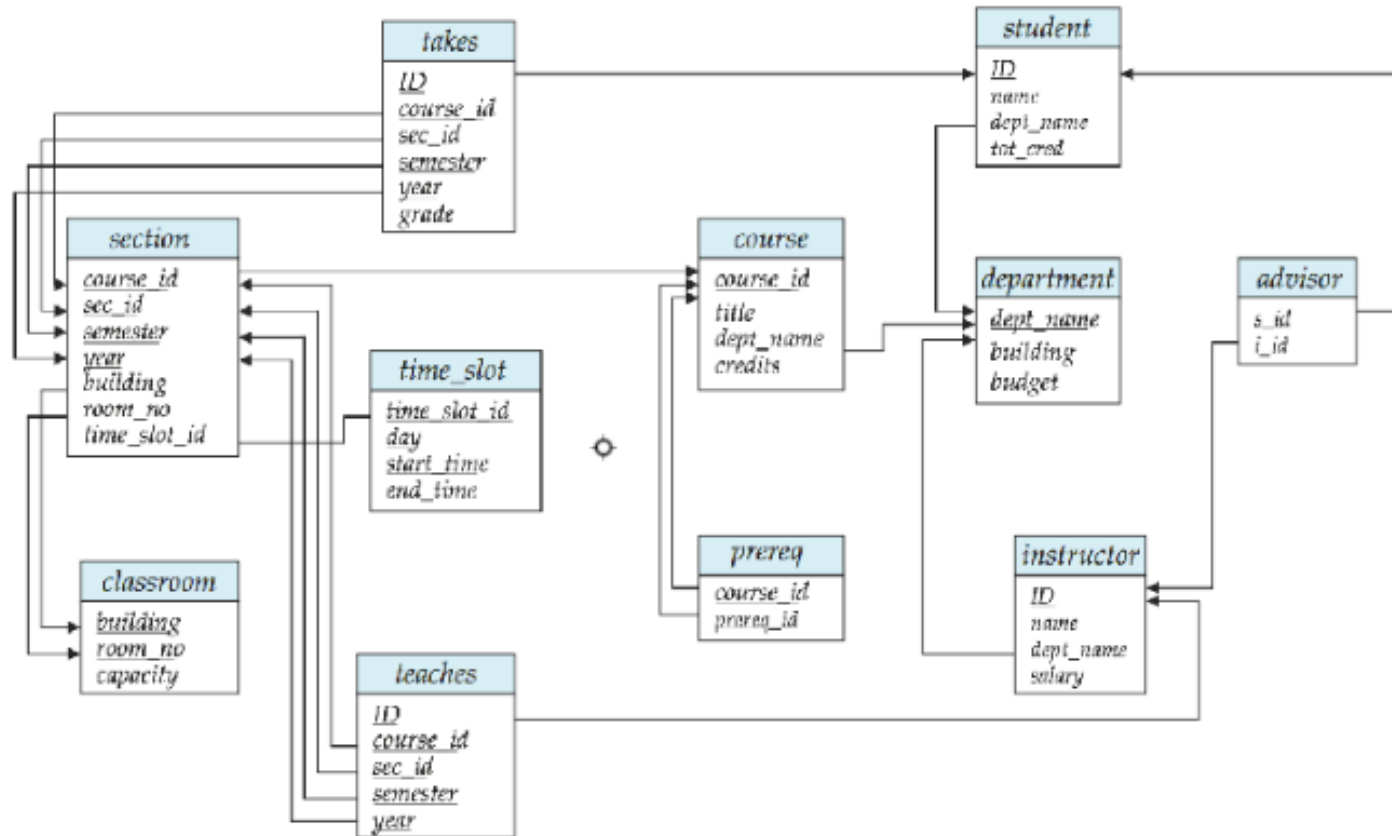
Multiple Tables

CMPSC 305 – Database Systems



ALLEGHENY COLLEGE

How to Connect Information? - No free-standing tables allowed!



The Basic Query Structure

The SQL data-manipulation language (DML) provides the ability to query information, and insert, delete and update tuples

A typical SQL code has the form:

```
SELECT A1, A2, ..., An  
FROM r1, r2, ..., rm  
WHERE P;
```

- A_i represents an attribute
- R_i represents a relation
- P is a predicate
- The result of an SQL query is a relation

The SELECT Clause

The SELECT clause filters out particular data from a table.

- SQL allows duplicates in relations as well as in query results.
- The SELECT statement has many optional clauses:
 - **WHERE** specifies which rows to retrieve.
 - **GROUP BY** groups rows sharing a property so that an aggregate function can be applied to each group.
 - **HAVING** selects among the groups defined by the GROUP BY clause.
 - **ORDER BY** specifies an order in which to return the rows.
 - **AS** provides an alias which can be used to temporarily rename tables or columns.

Given table 'T'

Table "T"	Query	Result												
<table><tr><th>C1</th><th>C2</th></tr><tr><td>1</td><td>a</td></tr><tr><td>2</td><td>b</td></tr></table>	C1	C2	1	a	2	b	<pre>SELECT * FROM T;</pre>	<table><tr><th>C1</th><th>C2</th></tr><tr><td>1</td><td>a</td></tr><tr><td>2</td><td>b</td></tr></table>	C1	C2	1	a	2	b
C1	C2													
1	a													
2	b													
C1	C2													
1	a													
2	b													
<table><tr><th>C1</th><th>C2</th></tr><tr><td>1</td><td>a</td></tr><tr><td>2</td><td>b</td></tr></table>	C1	C2	1	a	2	b	<pre>SELECT C1 FROM T;</pre>	<table><tr><th>C1</th></tr><tr><td>1</td></tr><tr><td>2</td></tr></table>	C1	1	2			
C1	C2													
1	a													
2	b													
C1														
1														
2														
<table><tr><th>C1</th><th>C2</th></tr><tr><td>1</td><td>a</td></tr><tr><td>2</td><td>b</td></tr></table>	C1	C2	1	a	2	b	<pre>SELECT * FROM T WHERE C1 = 1;</pre>	<table><tr><th>C1</th><th>C2</th></tr><tr><td>1</td><td>a</td></tr></table>	C1	C2	1	a		
C1	C2													
1	a													
2	b													
C1	C2													
1	a													
<table><tr><th>C1</th><th>C2</th></tr><tr><td>1</td><td>a</td></tr><tr><td>2</td><td>b</td></tr></table>	C1	C2	1	a	2	b	<pre>SELECT * FROM T ORDER BY C1 DESC;</pre>	<table><tr><th>C1</th><th>C2</th></tr><tr><td>2</td><td>b</td></tr><tr><td>1</td><td>a</td></tr></table>	C1	C2	2	b	1	a
C1	C2													
1	a													
2	b													
C1	C2													
2	b													
1	a													

The TeaDB

Build file: sandbox/teaDB/teaDB Build.txt

- `cat builder_teaTableDB.txt | sqlite3 teaDB.sqlite3`

(or just copy in the text into SQLite3!)

The SELECT Clause

Find everything in the Department table.

- `SELECT * FROM Department;`

Find all entries for dept's of the Department table

- `SELECT dept from Department;`

Count entries of dept's in Department table,

- `SELECT COUNT(dept) FROM department;`

The SELECT Clause

Find all unique entries for depts in Department table,

- `SELECT DISTINCT(dept) FROM department;`

Count unique entries of depts in Department table,

- `SELECT COUNT(DISTINCT(dept)) FROM Department;`

Return an exhaustive set of sandwiches that are being ordered.

- `SELECT DISTINCT(sandwich) FROM Tea;`

What query to use to count these types of sandwiches?

The WHERE Clause

The WHERE clause: conditions that the result must satisfy

- Corresponds to the selection predicate of the relational algebra
- Comparison results can be combined using the logical connectives and, or, and not
- Comparisons can be applied to results of arithmetic expressions

The WHERE Clause

Find out who is ordering a sandwich less than \$15 (from the new cost column)

- `SELECT * FROM tea WHERE cost < 15;`

Find out what kinds of sandwiches are going to each dept

- `SELECT department.dept, tea.sandwich FROM department, tea
WHERE department.id == tea.id;`

The WHERE Clause

Find department, Session material, sandwich type for orders of sandwiches less than \$15.

```
SELECT
    Department.id, Session.session,
    tea.sandwich, tea.cost
FROM
    Tea, Department, Session
WHERE
    cost < 15
AND
    Department.id == Session.id
AND
    Department.id == Tea.id;
```

The WHERE Clause

Find out which professors are presenting posters

- `SELECT * FROM session WHERE material == "poster";`
- `SELECT ID, material FROM session WHERE material == "poster";`

Find how who is presenting a poster, having what kind of sandwich which costs over \$10

- `SELECT session.ID, session.material, tea.sandwich, tea.cost FROM session, tea WHERE session.material == "poster" AND tea.cost >10 AND session.id == tea.id;`

New Database Tables!

```
cat campusDB_build.txt | sqlite3 CampusDB.sqlite3
```

Abbreviations in Queries

Find which students are working with what instructors.

- `SELECT Instructor.ID, Instructor.name, Instructor.studentId, Student.name, Student.Id FROM Instructor, Student WHERE Instructor.studentId == Student.ID;`

Shorter way to write query by using abbreviations

- `SELECT i.ID, i.name, i.studentId, s.name, s.Id FROM Instructor i, Student s WHERE i.studentId == s.ID;`

- The “**Instructor**” table name can be replaced with an **i**
- The “Student” table name can be replaced by an “s”

Aggregate Functions

These functions operate on the multiset of values of a column of a relation, and return a value

- avg: average value
- min: minimum value
- max: maximum value
- sum: sum of values
- count: number of values

Mathematical Functions

To find all instructors in Comp. Sci. dept with salary > 80000

- ```
SELECT
 name FROM instructor
WHERE
 department = "CompSci"
AND
 salary > 80000;
```



# Mathematical Functions

## Using built-in functions

- `SELECT AVG (salary) FROM instructor WHERE deptName = "CompSci";`
- `SELECT MIN (salary) FROM instructor WHERE deptName = "CompSci";`
- `SELECT MAX (salary) FROM instructor WHERE deptName = "CompSci";`

# Mathematical Functions

To find all instructors in Comp. Sci. dept with salary > 80000

- `SELECT name FROM instructor  
WHERE deptName = "CompSci" AND salary > 80000;`

Using functions

- `SELECT SUM (salary) FROM instructor WHERE  
deptName = "CompSci";`
- `SELECT COUNT (salary) FROM instructor WHERE  
deptName = "CompSci";`

## Attention to the WHERE clause

Find the ID, name and total credit students who are taking a course where the total credit is 3 or 4 hours.

Why will “AND” NOT work

- `SELECT ID, name, totCred FROM student WHERE totCred == "3" OR totCred == "4";`

## Attention to the WHERE clause

Watch out for cross products that give no usable information!!

- `SELECT s.name, i.name from student s, instructor i WHERE s.deptName == i.deptName and s.deptName == "CompSci";`
- Common Solution – Use two queries instead
  - `SELECT s.name from student s WHERE s.deptName == "CompSci";`
  - `SELECT i.name from instructor i WHERE i.deptName == "CompSci";`

## Using Count and Count(Distinct(...))

Find the number of tuples in the course relation

- `SELECT COUNT(credits) FROM course;`
- `SELECT COUNT(distinct(credits)) FROM course;`
- `SELECT COUNT (*) FROM course;`
- `SELECT COUNT(distinct(*)) FROM course;`
- Question: Why will the above distinct line not work?

# Removing Tables or Data

- `DROP TABLE IF EXISTS student`  
Deletes the table student and its contents if present
- `DROP TABLE student`  
Deletes the table student and its contents, report error if table not present
- `DELETE FROM student`  
Deletes all contents from table student, but retains table

Play with your database!

Remember, you can use your builder file to re-create the database if it becomes corrupt or unstable.

# Changing Table Contents

## ALTER TABLE

- `Alter TABLE r ADD AD`
- where A is the name of the attribute to be added to relation r and D is the domain of A.
- All tuples in the relation are assigned null as the value for the new attribute.
- EX: `ALTER TABLE Department ADD Email varchar;`

## Change name of table:

- `ALTER TABLE Department RENAME TO newDept;`

# Changing Table Contents

## Add a column to a table

- ALTER TABLE course ADD COLUMN courseTag VARCHAR;
- Check your additional column:  
    .schema course

## Dropping a column of a table

- ALTER TABLE course DROP COLUMN courseTag;
- Check your additional column:  
    .schema course



# Complex Queries

Instructor names, IDs and their Students?

```
SELECT
 instructor.name, instructor.id,
 instructor.studentID, student.ID,
 student.name
FROM
 instructor, student
WHERE
 student.id == instructor.studentID;
```

# Complex Queries

## Output

Miller|10101|S1|S1|Michaels  
Johnson|10102|S1|S1|Michaels  
Charleson|10103|S2|S2|Peterson  
Thompson|10104|S2|S2|Peterson  
Mauler|10105|S3|S3|Mullen  
...  
Farber|10112|S5|S5|Beuller

## Try these with campusDB.sqlite3!

- What is the average salary of computer science teachers?
- What is the average salary of computer science teachers who make less than \$98000?



# Try these with campusDB.sqlite3!

- What are the average salaries of instructors who worked during the Spring?
- What are the average salaries of instructors who worked during the Fall?



## Try these with campusDB.sqlite3!

- What are the Instructor names and their IDs who taught which Students (show names and IDs)?
- What are the Instructor IDs their Student's IDs in cases where the instructors and students are NOT in the same department?



## Try these with campusDB.sqlite3!

- What are the Instructor names and their IDs who taught which Students (show names and IDs) for classes taught in the year 2010?
- What are the Instructor names and their IDs who taught which Students (show names and IDs) for classes taught in the year 2010. In which semester were they teaching?
- Come-up with your own complex question and query solution.



**THINK**