

# Advanced queries, joins and aggregates

CMPSC 305 – Database Systems



ALLEGHENY COLLEGE

## Joins: Bringing Data Together



- The SQLite3 join-clause is used to combine records from two or more tables in a database.
- A JOIN is a means for combining fields from two tables by using values common to each.

# Joins: Visual Definitions

As Venn Diagrams

Inner Join




Table 1

Table 2

Left Join




Table 1

Table 2

Right Join




Table 1

Table 2

Full Outer Join





Table 1


Table 2

# SQL Code and Venn Diagrams


## SQL JOINS




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

# Joins: Visual Definitions


## Combining Tables

Left




key	val_x	val_y
1	x1	y1
2	x2	y2
3	x3	NA

Right



key	val_x	val_y
1	x1	y1
2	x2	y2
4	NA	y3

Full



key	val_x	val_y
1	x1	y1
2	x2	y2
3	x3	NA
4	NA	y3

# An Explanation of Terms

## SQL joins

- An inner join will return records that have matching values in both tables.
- A left outer join will return all records from the left table and the matched records from the right table.
- A right outer join will return all records from the right table and the matched records from the left table.
- A full outer join will return all records when there is a match from either table.


## Inner Joins

Student ID	Name
1001	A
1002	B
1003	C
1004	D

Table A

Student ID	Department
1004	Mathematics
1005	Mathematics
1006	History
1007	Physics
1008	Computer Science

Table B



Student ID	Name	Department
1004	D	Mathematics

# Inner Joins

File: /sandbox/fruitJoin.txt

## Create two tables

```
DROP TABLE IF EXISTS TableA;  
CREATE TABLE TableA (  
  fruit VARCHAR,  
  colour VARCHAR);
```

```
DROP TABLE IF EXISTS TableB;  
CREATE TABLE TableB (  
  fruit VARCHAR,  
  colour VARCHAR);
```



# Inner joins

File: /sandbox/fruitJoin.txt

## Populate the tables

```
INSERT INTO TableA VALUES ("Lemons_A","Yellow");  
INSERT INTO TableA VALUES ("Apples_A","Red");  
INSERT INTO TableA VALUES ("Grapes_A","Purple");
```

```
INSERT INTO TableB VALUES ("Lemons_B","Yellow");  
INSERT INTO TableB VALUES ("Apples_B","Red");  
INSERT INTO TableB VALUES ("Oranges_B", "Orange");  
INSERT INTO TableB VALUES ("Grapes_B","Purple");
```

# Inner joins

File: /sandbox/fruit- innerJoin.txt

Use INNER JOIN to query

```
.tables
```

```
SELECT * from TableA;
```

```
SELECT* from TableB;
```

```
SELECT
```

```
    TableA.fruit,  
    TableA.colour,  
    TableB.colour,  
    TableB.fruit
```

```
FROM
```

```
    TableA
```

```
INNER JOIN
```

```
    TableB ON TableB.colour == TableA.colour;
```

## Inner joins

### Output

```
Lemons_A | Yellow | Yellow | Lemons_B  
Apples_A  | Red    | Red    | Apples_B  
Grapes_A  | Purple | Purple | Grapes_B
```

## Where vs Inner Joins

```
SELECT
    TableA.fruit,
    TableA.colour,
    TableB.colour,
    TableB.fruit
FROM
    TableA, TableB
WhERE
    TableB.colour = TableA.colour;
```

# Left Join

Matches entries from LEFT table to the other table

## Setup Tables

```
DROP TABLE IF EXISTS Employees;
```

```
CREATE TABLE Employees (  
    EmployeeID INT PRIMARY KEY, LastName VARCHAR, DepartmentID INT,  
    FirstName VARCHAR  
);
```

```
DROP TABLE IF EXISTS Departments;
```

```
CREATE TABLE Departments (  
    DepartmentID INT PRIMARY KEY, DepartmentName VARCHAR  
);
```

## Left Join

```
INSERT INTO Employees (  
    EmployeeID, FirstName,  
    LastName, DepartmentID)
```

```
VALUES
```

```
    (1, 'John', 'Doe', 1),  
    (2, 'Jane', 'Smith', 2),  
    (3, 'Bob', 'Johnson', 1),  
    (4, 'Alice', 'Williams', NULL);
```

```
INSERT INTO Departments (DepartmentID, DepartmentName)
```

```
VALUES
```

```
    (1, 'HR'),  
    (2, 'IT'),  
    (3, 'Finance');
```

## Left Join

```
/*Perform a LEFT JOIN to retrieve a list of all employees and their departments*/  
SELECT  
    e.EmployeeID, e.FirstName, e.LastName, d.DepartmentName  
FROM  
    Employees e  
LEFT JOIN  
    Departments d  
ON  
    e.DepartmentID = d.DepartmentID;
```

EmployeeID	FirstName	LastName	DepartmentName
------------	-----------	----------	----------------

1	John	Doe	HR
2	Jane	Smith	IT
3	Bob	Johnson	HR
4	Alice	Williams	NULL

# Right Join

Matches entries from RIGHT table to the other table

```
/*Perform a RIGHT JOIN to retrieve a list of all departments, even if they have no employees.*/  
SELECT  
    e.EmployeeID, e.FirstName, e.LastName, d.DepartmentName  
FROM  
    Employees e  
RIGHT JOIN  
    Departments d  
ON  
    e.DepartmentID = d.DepartmentID;
```

EmployeeID	FirstName	LastName	DepartmentName
------------	-----------	----------	----------------

1	John	Doe	HR
2	Jane	Smith	IT
3	Bob	Johnson	HR
NULL	NULL	NULL	Finance



# Cross joins

Cross product derived from both tables

```
DROP TABLE IF EXISTS ranks;
CREATE TABLE ranks (
    rank TEXT NOT NULL
);
DROP TABLE IF EXISTS suits;
CREATE TABLE suits (
    suit TEXT NOT NULL
);

INSERT INTO ranks(rank)
VALUES('2'),('3'),('4'),('5'),('6'),('7'),('8'),('9'),('10'),('J'),('Q'),('K'),('A');

INSERT INTO suits(suit) VALUES('Clubs'),('Diamonds'),('Hearts'),('Spades');

SELECT rank, suit
      FROM ranks
     CROSS JOIN suits
 ORDER BY suit;
```

# Cross joins: All Card Pairs

Cross join: output

2 | Clubs

...

J | Clubs




Q | Clubs

K | Clubs

A | Clubs

...

A | Spades



## New Database



(A New Database!)

# New Database

Schema: Red boxes are the tables of today's database study

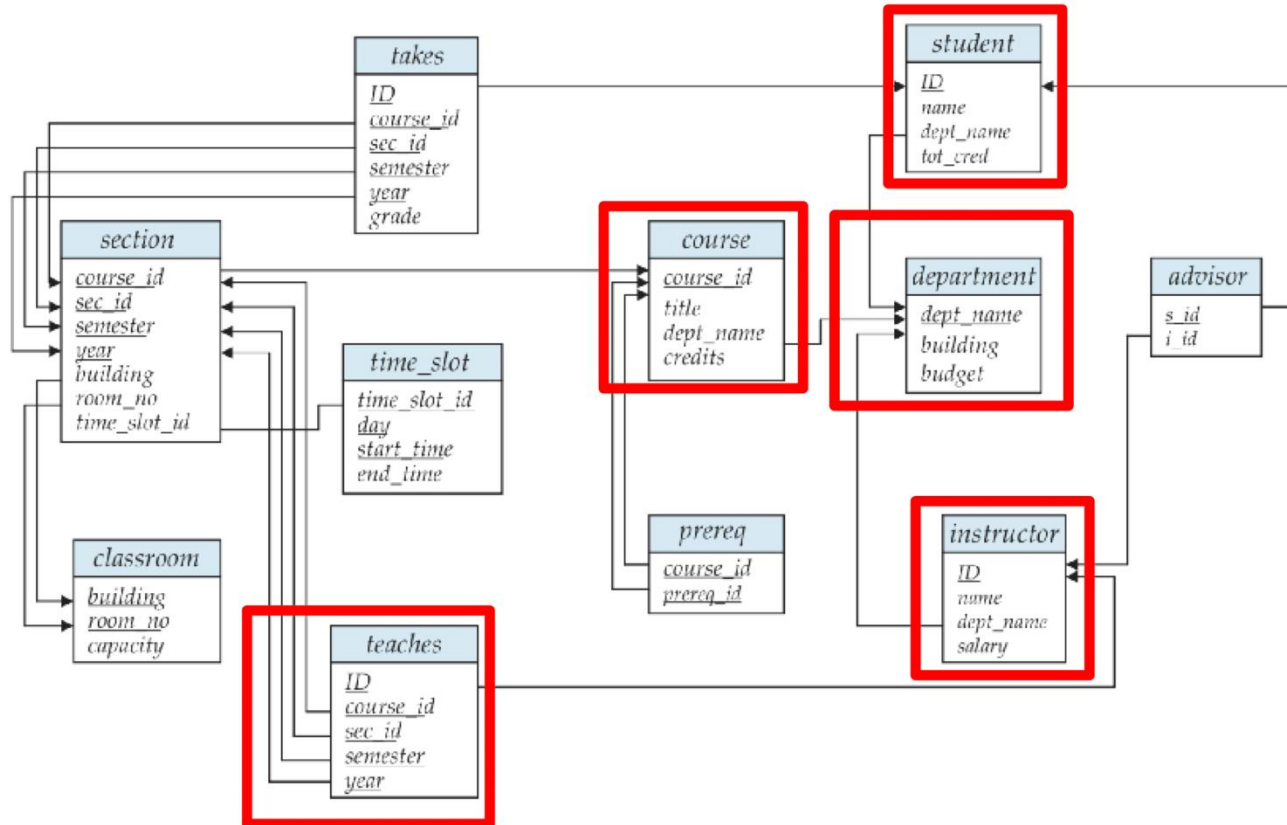


Figure 2.8 Schema diagram for the university database.

## New Database

- Find the database maker file, campusDB build.txt, in your sandbox directory

```
cat campusDB_build.txt | sqlite3 myCampusDB.sqlite3
```

# Set Operations

## OR & AND

- OR: Find all deptNames in the **UNION** of Instructor and Course
- select deptName from Instructor **UNION** select deptName from course;
- select distinct(deptName) from Instructor;
- AND: Find all deptNames in the **INTERSECT** of Instructor and Course
- select deptName from Instructor **INTERSECT** select deptName from Course;
- select distinct(Instructor.deptName) from Instructor, Course  
where Instructor.deptName == Course.deptName;

## Set Operations

- `select distinct(deptName) from Instructor;`
  - `select distinct(deptName) from Course;`
- 
- The **EXCEPT** operator compares the result sets of two queries and returns distinct rows from the left query that are not in the output by the right query.
  - Find all deptNames different to both the Instructor and Course
  - Check these two queries below. Why is the output different?
- 
- `select deptName from Instructor EXCEPT select deptName from Course;`
  - `select deptName from Course EXCEPT select deptName from Instructor;`

## AS clauses

- The AS clause is used to rename relations; useful for reducing necessary code in queries
- Ex: For all instructors in the university who have taught some course, find their names and the course ID of all their taught courses

```
Select I.name, T.courseID  
FROM Instructor AS I, Teaches AS T  
WHERE I.ID= T.ID;
```

- On the second line:
- the Instructor table is renamed to I
- the Teaches table is renamed to T.



## AS clauses

- Another reason to rename a relation is a case where we wish to compare tuples in the same relation.
- We then need to take the Cartesian product of a relation with itself and, without renaming, it becomes impossible to distinguish one tuple from the other.
- Suppose that we want to write the query, find the names of all instructors whose salary is greater than at least one instructor in the Math department.
- ```
SELECT DISTINCT T.name
FROM Instructor as T,
Instructor AS S
WHERE T.salary > S.salary and S.deptName == "Math";
```

## AS clauses

- Find all names of common teachers in Instructor and Teaches tables

Use AS to implement variables attributes to hold places

```
select distinct(Instructor.name) as newName  
From Instructor, teaches  
Where Instructor.ID = teaches.ID and newName == "Thompson";
```

## AS clauses

- Find the names of all Instructors whose salary is greater than at least one Instructor in the Math department.
- `select distinct(T.name) from Instructor as T,  
Instructor as S where T.salary > S.salary and  
S.deptName == "Math";`
- `select distinct T.name, T.salary from Instructor as T, Instructor as S  
where T.salary > S.salary and S.deptName == "Math";`
- Reference: `select * from Instructor;`