

Intro to MongoDB

CMPSC 305 – Database Systems



ALLEGHENY COLLEGE

The Problem with SQL

firstName	lastName	primaryAddr

- Let's say that we have a (perfectly) working SQL table
- The schema has been designed and coded for current data requirements

Table Update (i)

firstName	lastName	primaryAddr	secondAddr


- The data we collect has changed.
- We need to update our schema for the new data requirements

Table Update (ii)

firstName	lastName	primaryAddr	secondAddr	thirdAddr

- Our needs have changed again, and the SQL table must be updated.
- The schema is reprogrammed

Expectations

firstName	lastName	primaryAddr	secondAddr	thirdAddr
				
				
				
				

- We expect that the table will be full when in use
- Expectations are not always fulfilled...

In Reality, Much Data is Missing


firstName	lastName	primaryAddr	secondAddr	thirdAddr
Smiley Face (Blue)				
Smiley Face (Teal)				
Smiley Face (Red)				

- But, in reality, much of the table is empty!
- The table can easily get huge and be hard to manage.

We Might Ask Ourselves...

- What can we do to stop having to redesign our database schema with our changing data?
- Is SQL the right type of database management system for our changing data requirements?

A NoSQL Database Management System



mongoDB®

- NoSQL: Not Only SQL database systems that support SQL-like query languages, but are used increasingly in big data applications and real-time web applications.
 - The stored data is allowed to change
-
- <https://www.mongodb.com/>

Philosophy of MongoDB

Non-relational DB

- Document Identifiers (id) will be created for each document, field name reserved by system
- Application tracks the schema and mapping
- Uses JSON, BSON (B for binary inputs)
- Written in C++
- Supports APIs (drivers) in many computer languages

JavaScript, Python, Golang, Ruby, Perl, Java, Java Scala, C#,
C++, Haskell, Erlang

Database Language Guide


SQL systems versus NoSQL

RDBMS		MongoDB
Database	➡	Database
Table, View	➡	Collection
Row	➡	Document (BSON)
Column	➡	Field
Index	➡	Index
Join	➡	Embedded Document
Foreign Key	➡	Reference
Partition	➡	Shard

- The terms are different, but their meanings are similar
- Schema-less, collections (like tables) are populated by any data
- Documents are similar to the tuples of Sqlite3 programming

Database Language Guide

SQL systems versus NoSQL




Collection

- No pre-defined data schema
 - Data may be entered in absence of a defined schema
- Documents (rows) of collections (DB's) may have different types of data

Schema Free

Mostly similar documents



- Sometimes not all the data is available to create a document.
- The query interprets missing data as NULL entries

Styles of Storing Data

Relational Database

Student_Id	Student_Name	Age	College
1001	Chaitanya	30	Beginnersbook
1002	Steve	29	Beginnersbook
1003	Negan	28	Beginnersbook



MongoDB

```
{  
  "_id": ObjectId("....."),  
  "Student_Id": 1001,  
  "Student_Name": "Chaitanya",  
  "Age": 30,  
  "College": "Beginnersbook"  
}  
  
{  
  "_id": ObjectId("....."),  
  "Student_Id": 1002,  
  "Student_Name": "Steve",  
  "Age": 29,  
  "College": "Beginnersbook"  
}  
  
{  
  "_id": ObjectId("....."),  
  "Student_Id": 1003,  
  "Student_Name": "Negan",  
  "Age": 28,  
  "College": "Beginnersbook"  
}
```

JSON and MongoDB Code

- **Data is in name / value pairs**
 - A name/value pair consists of a field name followed by a colon, followed by a value:
Example: { "name": "R2-D2" }
- **Data is separated by commas**
 - Example: { "name": "R2-D2", race : "Droid"}
- **Curly braces hold objects**
 - Example: {"name": "R2-D2", race : "Droid", affiliation: "rebels"}
- **An array is stored in brackets []**
 - Example [{ "name": "R2-D2", race : "Droid", affiliation: "rebels"}, { "name": "Yoda", affiliation: "rebels"}]

CRUD Operations

Create, Read, Update and Delete

- **Db.collection** specifies the collection or the table in which to store the document (tuple)
- **Create**
 - db.collection.insert()
 - db.collection.save()
 - db.collection.update()
- **Read**
 - db.collection.find()
 - db.collection.findOne()
- **Update**
 - db.collection.update()
- **Delete**
 - db.collection.remove()

Let's code

Online coding to test the new DB



Let's code

Online MongoDB

The screenshot shows the myCompiler.io interface for MongoDB. The URL in the browser bar is `mycompiler.io/new/mongodb`. The interface includes a search bar for programs, language selection (English), and user profile information. A title input field says "Enter a title...", and a tags input field says "+ Add tags (upto 10)". Below the code editor, there are dropdowns for "MongoDB" and "Run/Save" buttons. The code editor contains the following MongoDB script:

```
1 db.students.insertMany([
2   { id: 1, name: 'Ryan', gender: 'M' },
3   { id: 2, name: 'Joanna', gender: 'F' }
4 ]);
5 db.students.find({ gender: 'F' });
```

The output panel is currently empty and displays the placeholder text "(Run the program to view its output)".

- <https://www.mycompiler.io/new/mongodb>

Let's code

output

Output

```
mycompiler_mongodb> ... . . . . {
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('6906c2e5e9b991bbe56b128c'),
    '1': ObjectId('6906c2e5e9b991bbe56b128d')
  }
}
mycompiler_mongodb> [
{
  _id: ObjectId('6906c2e5e9b991bbe56b128d'),
  id: 2,
  name: 'Joanna',
  gender: 'F'
}
]
mycompiler_mongodb>
```

[Execution complete with exit code 0]

Simple Collections

Enter data as JSON code

Insert each document individually into the Furniture collection

```
db.Furniture.drop()  
    db.Furniture.insertOne({chair:"wood"})  
    db.Furniture.insertOne({chair:"metal"})  
    db.Furniture.insertOne({chair:"plastic"})  
    db.Furniture.insertOne({table:"glass"})  
    db.Furniture.insertOne({table:"wood"})  
    db.Furniture.insertOne({table:"metal"})  
    db.Furniture.insertOne({lamp:"brass"})  
    db.Furniture.insertOne({lamp:"glass"})  
    db.Furniture.insertOne({lamp:"silver"})
```

```
db.Furniture.find({searchSpace},{showAttrib:1})
```

```
Find everything: db.Furniture.find({},{}))
```

Let's code

output



MongoDB ▾



```
1 db.Furniture.drop()  
2 db.Furniture.insertOne({chair:"wood"})  
3 db.Furniture.insertOne({chair:"metal"})  
4 db.Furniture.insertOne({chair:"plastic"})  
5 db.Furniture.insertOne({table:"glass"})  
6 db.Furniture.insertOne({table:"wood"})  
7 db.Furniture.insertOne({table:"metal"})  
8 db.Furniture.insertOne({lamp:"brass"})  
9 db.Furniture.insertOne({lamp:"glass"})  
10 db.Furniture.insertOne({lamp:"silver"})  
11  
12 db.Furniture.find({}, {lamp:1})  
13
```

Find all lamps: db.Furniture.find({}, {lamp:1})

Let's code

output

```
mycompiler_mongodb>
mycompiler_mongodb> [
  { _id: ObjectId('6906c3d68ba14fc50e6b128c') },
  { _id: ObjectId('6906c3d68ba14fc50e6b128d') },
  { _id: ObjectId('6906c3d78ba14fc50e6b128e') },
  { _id: ObjectId('6906c3d78ba14fc50e6b128f') },
  { _id: ObjectId('6906c3d78ba14fc50e6b1290') },
  { _id: ObjectId('6906c3d78ba14fc50e6b1291') },
  { _id: ObjectId('6906c3d78ba14fc50e6b1292'), lamp: 'brass' },
  { _id: ObjectId('6906c3d78ba14fc50e6b1293'), lamp: 'glass' },
  { _id: ObjectId('6906c3d78ba14fc50e6b1294'), lamp: 'silver' }
]
mycompiler_mongodb>
```

[Execution complete with exit code 0]

Find all lamps: db.Furniture.find({}, {lamp:1})

Let's code

output



MongoDB ▾



```
1 db.Furniture.drop()  
2 db.Furniture.insertOne({chair:"wood"})  
3 db.Furniture.insertOne({chair:"metal"})  
4 db.Furniture.insertOne({chair:"plastic"})  
5 db.Furniture.insertOne({table:"glass"})  
6 db.Furniture.insertOne({table:"wood"})  
7 db.Furniture.insertOne({table:"metal"})  
8 db.Furniture.insertOne({lamp:"brass"})  
9 db.Furniture.insertOne({lamp:"glass"})  
10 db.Furniture.insertOne({lamp:"silver"})  
11  
12 db.Furniture.find({}, {lamp:1}).pretty()  
13
```

Find all lamps, use formatting: db.Furniture.find({}, {lamp:1}).pretty()

Let's code

output

```
mycompiler_mongodb>
mycompiler_mongodb> [
  { _id: ObjectId('6906c42393cf764426b128c') },
  { _id: ObjectId('6906c42493cf764426b128d') },
  { _id: ObjectId('6906c42493cf764426b128e') },
  { _id: ObjectId('6906c42493cf764426b128f') },
  { _id: ObjectId('6906c42493cf764426b1290') },
  { _id: ObjectId('6906c42493cf764426b1291') },
  { _id: ObjectId('6906c42493cf764426b1292'), lamp: 'brass' },
  { _id: ObjectId('6906c42493cf764426b1293'), lamp: 'glass' },
  { _id: ObjectId('6906c42493cf764426b1294'), lamp: 'silver' }
]
mycompiler_mongodb>
```

[Execution complete with exit code 0]

Find all lamps, use formatting: db.Furniture.find({}, {lamp:1}).pretty()

Simple Collections

Simple Example of Queries

Query all documents in the Furniture collection

```
db.Furniture.find({},{})
```

Query all Lamp types across all collections

```
db.Furniture.find({}, {lamp:1})
```

```
db.Furniture.find({}, {lamp:1, _id:0})
```

Simple Collections

Simple Example of Queries

Query Lamp types from the Furniture collection

```
// SELECT lamp FROM Furniture WHERE lamp == ``brass'';  
db.Furniture.find({lamp:"brass"})  
db.Furniture.find({lamp:"glass"})  
db.Furniture.find({lamp:"silver"})  
  
// do not show object id's  
db.Furniture.find({lamp:"silver"}, {_id:0})
```

Insert many documents into the Inventory collection

Inserting

```
db.inventory.insertMany([  
  { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },  
  { item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A" },  
  { item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D" },  
  { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },  
  { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" }  
]);
```

Insert many documents into the Inventory collection

```
SELECT * FROM inventory
```

```
db.inventory.find( {},{} )  
db.inventory.find( {},{} ).pretty()
```

```
SELECT item FROM inventory
```

```
db.inventory.find({},{ "item":1 }).pretty()
```

```
SELECT * FROM inventory WHERE item == “postcard”
```

```
db.inventory.find({ "item": "postcard"},{})  
db.inventory.find({ "item": "postcard"},{}).pretty()
```

Queries from the Inventory collection

SELECT * FROM inventory WHERE status = "D"

```
db.inventory.find( { status: "D" } )
```

SELECT * FROM inventory WHERE status in ("A", "D")

```
db.inventory.find({status:{ $in: [ "A", "D" ]}})
```

SELECT * FROM inventory WHERE status == "D")

```
db.inventory.find({ status: "D" },{})
```

Show me where the size = "h" and size = 10)

```
db.inventory.find( {"size.h":10} ).pretty()
```

See more on this at <https://www.mongodb.com/docs/manual/tutorial/query-documents/>

Consider this!



THINK

- Can you go back to the above examples to query other fascinating information?
- Can you create and populate a new Mongo database?
- Can you write sophisticated queries in your database?